

This article is showing a geometric and intuitive explanation of the covariance matrix and the way it describes the shape of a data set. We will describe the geometric relationship of the covariance matrix with the use of linear transformations and eigendecomposition.

## Introduction

Before we get started, we shall take a quick look at the difference between covariance and variance. **Variance measures the variation of a single random variable (like the height of a person in a population), whereas covariance is a measure of how much two random variables vary together (like the height of a person and the weight of a person in a population).** The formula for variance is given by

$$\sigma_x^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

where  $n$  is the number of samples (e.g. the number of people) and  $\bar{x}$  is the mean of the random variable  $x$  (represented as a vector). The covariance  $\sigma(x, y)$  of two random variables  $x$  and  $y$  is given by

$$\sigma(x, y) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

with  $n$  samples. The variance  $\sigma_x^2$  of a random variable  $x$  can be also expressed as the covariance with itself by  $\sigma(x, x)$ .

## Covariance Matrix

With the covariance we can calculate entries of the covariance matrix, which is a square matrix given by  $C_{i,j} = \sigma(x_i, x_j)$  where  $C \in \mathbb{R}^{d \times d}$  and  $d$  describes the dimension or number of random variables of the data (e.g. the number of features like height, width, weight, ...). Also the covariance matrix is symmetric since  $\sigma(x_i, x_j) = \sigma(x_j, x_i)$ . The diagonal entries of the covariance matrix are the variances and the other entries are the covariances. For this reason, the covariance matrix is sometimes called the `_variance-covariance matrix_`. The calculation for the covariance matrix can be also expressed as

$$C = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(X_i - \bar{X})^T$$

	$x$	$y$
$x$	$\sigma_x^2$	$\sigma(x,y)$
$y$	$\sigma(y,x)$	$\sigma_y^2$

where our data set is expressed by the matrix  $X \in \mathbb{R}^{n \times d}$ . Following from this equation, the covariance matrix can be computed for a data set with zero mean with  $C = \frac{XX^T}{n-1}$  by using the semi-definite matrix  $XX^T$ .

In this article, we will focus on the two-dimensional case, but it can be easily generalized to more dimensional data. Following from the previous equations the covariance matrix for two dimensions is given by

$$C = \begin{pmatrix} \sigma(x, x) & \sigma(x, y) \\ \sigma(y, x) & \sigma(y, y) \end{pmatrix}$$

We want to show how linear transformations affect the data set and in result the covariance matrix. First we will generate random points with mean values  $\bar{x}$ ,  $\bar{y}$  at the origin and unit variance  $\sigma_x^2 = \sigma_y^2 = 1$  which is also called [white noise](#) and has the identity matrix as the covariance matrix.

```
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

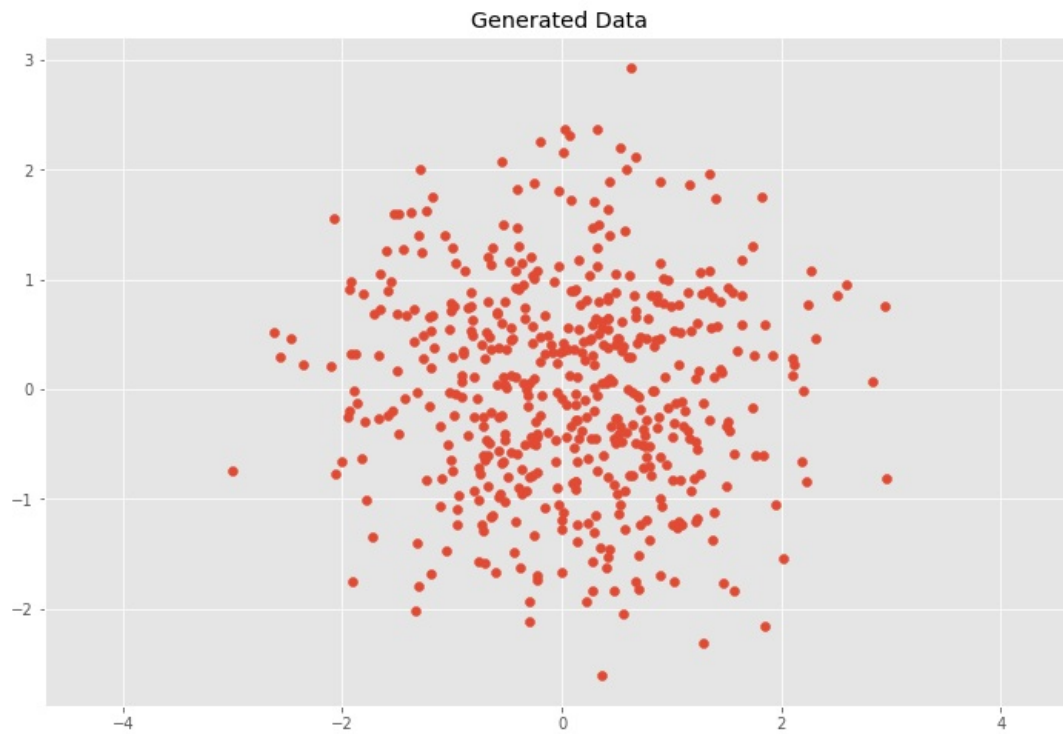
```
plt.style.use('ggplot')
plt.rcParams['figure.figsize'] = (12, 8)
```

```
# Normal distributed x and y vector with mean 0 and standard deviation 1
x = np.random.normal(0, 1, 500)
y = np.random.normal(0, 1, 500)
X = np.vstack((x, y)).T
```

```
plt.scatter(X[:, 0], X[:, 1])
plt.title('Generated Data')
plt.axis('equal');
Copy
```

$$\text{mean of } x : \frac{1}{n} \sum_{i=1}^n x_i$$

standard deviation measures the amount of variation or dispersion of a set of values  
low standard deviation  $\rightarrow$  values are closer to the mean



This case would mean that  $x$  and  $y$  are independent (or uncorrelated) and the covariance matrix  $C$  is

$$C = \begin{pmatrix} \sigma_x^2 & 0 \\ 0 & \sigma_y^2 \end{pmatrix}$$

We can check this by calculating the covariance matrix

```
# Covariance
def cov(x, y):
    xbar, ybar = x.mean(), y.mean()
    return np.sum((x - xbar)*(y - ybar))/(len(x) - 1)

# Covariance matrix
def cov_mat(X):
    return np.array([[cov(X[0], X[0]), cov(X[0], X[1]), \
                      [cov(X[1], X[0]), cov(X[1], X[1])]])

# Calculate covariance matrix
cov_mat(X.T) # (or with np.cov(X.T))
array([[ 1.008072 , -0.01495206],
       [-0.01495206,  0.92558318]])
```

Copy

Which approximately gives us our expected covariance matrix with variances  $\sigma_x^2 = \sigma_y^2 = 1$  .

## Linear Transformations of the Data Set

Next, we will look at how transformations affect our data and the covariance matrix  $C$  . We will transform our data with the following scaling matrix.

$$S = \begin{pmatrix} s_x & 0 \\ 0 & s_y \end{pmatrix}$$

where the transformation simply scales the  $x$  and  $y$  components by multiplying them by  $s_x$  and  $s_y$  respectively. What we expect is that the covariance matrix  $C$  of our transformed data set will simply be

$$C = \begin{pmatrix} (s_x\sigma_x)^2 & 0 \\ 0 & (s_y\sigma_y)^2 \end{pmatrix}$$

which means that we can extract the scaling matrix from our covariance matrix by calculating  $S = \sqrt{C}$  and the data is transformed by  $Y = SX$  .

```
# Center the matrix at the origin
X = X - np.mean(X, 0)
```

```
# Scaling matrix
sx, sy = 0.7, 3.4
```

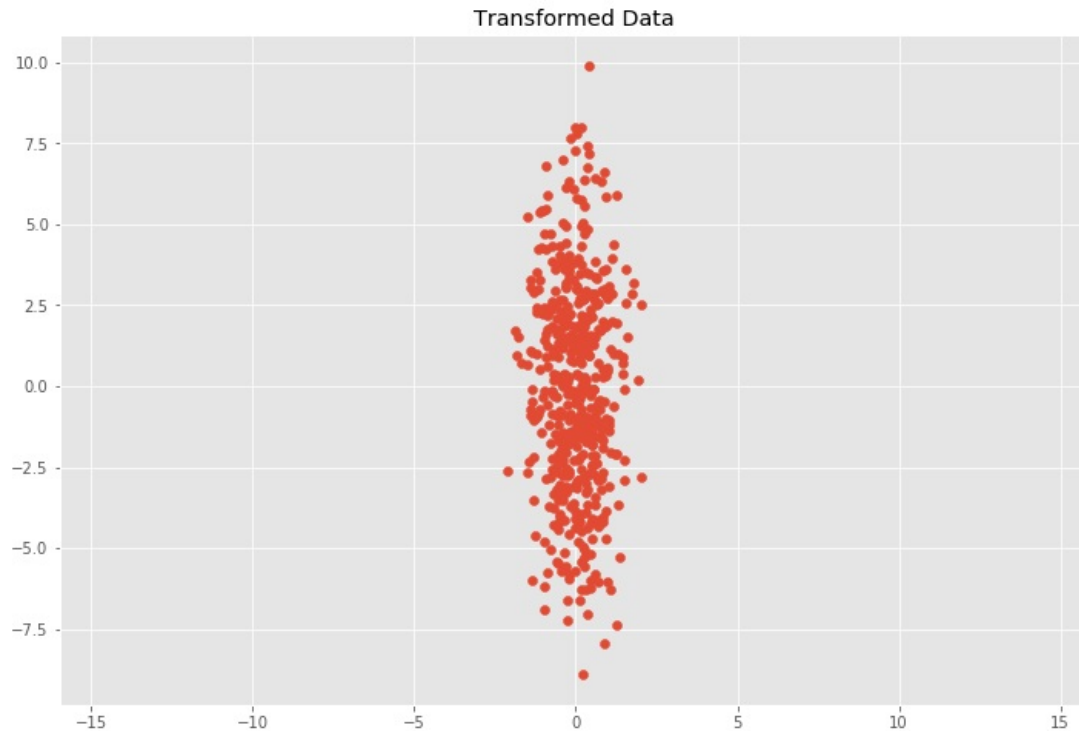
```
Scale = np.array([[sx, 0], [0, sy]])

# Apply scaling matrix to X
Y = X.dot(Scale)

plt.scatter(Y[:, 0], Y[:, 1])
plt.title('Transformed Data')
plt.axis('equal')

# Calculate covariance matrix
cov_mat(Y.T)
array([[ 0.50558298, -0.09532611],
       [-0.09532611, 10.43067155]])
```

Copy



We can see that this does in fact approximately match our expectation with  $0.7^2 = 0.49$  and  $3.4^2 = 11.56$  for  $(s_x\sigma_x)^2$  and  $(s_y\sigma_y)^2$ . This relation holds when the data is scaled in  $x$  and  $y$  direction, but it gets more involved for other linear transformations.

Now we will apply a linear transformation in the form of a transformation matrix  $T$  to the data set which will be composed of a two dimensional [rotation matrix](#)  $R$  and the previous scaling matrix  $S$  as follows

$$T = RS$$

where the rotation matrix  $R$  is given by

$$R = \begin{pmatrix} \cos(\theta) & -\sin(\theta) \\ \sin(\theta) & \cos(\theta) \end{pmatrix}$$

where  $\theta$  is the rotation angle. The transformed data is then calculated by  $Y = TX$  or  $Y = RSX$ .

```
# Scaling matrix
sx, sy = 0.7, 3.4
Scale = np.array([[sx, 0], [0, sy]])

# Rotation matrix
theta = 0.77*np.pi
c, s = np.cos(theta), np.sin(theta)
Rot = np.array([[c, -s], [s, c]])

# Transformation matrix
T = Scale.dot(Rot)

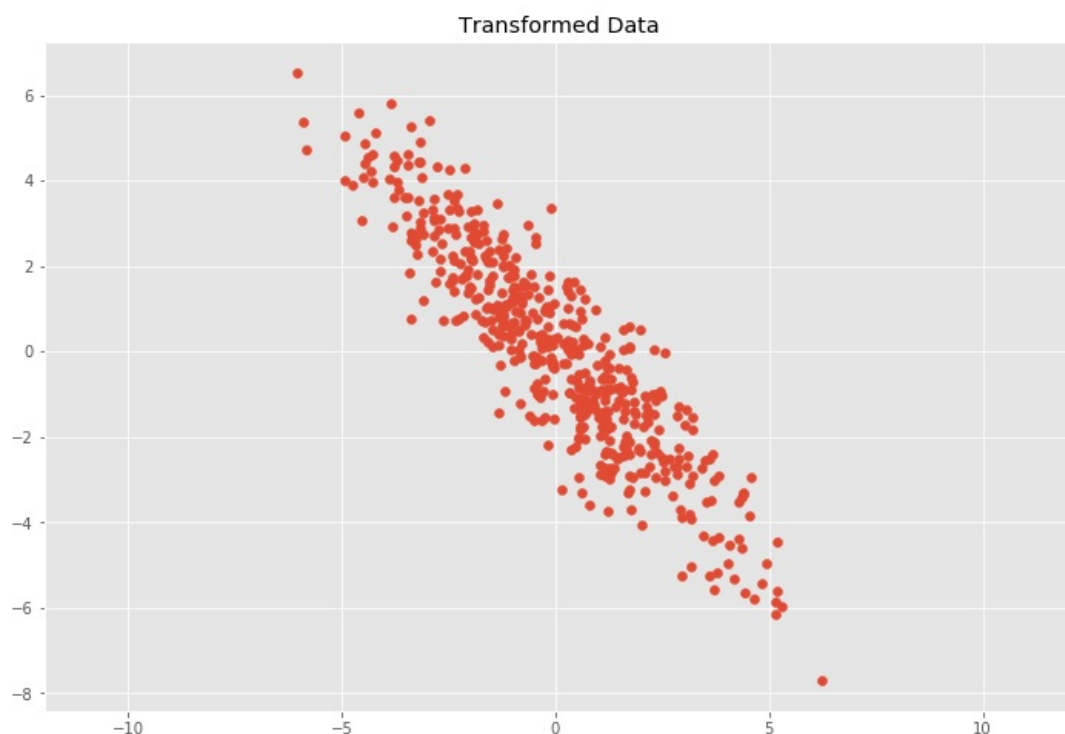
# Apply transformation matrix to X
Y = X.dot(T)

plt.scatter(Y[:, 0], Y[:, 1])
plt.title('Transformed Data')
plt.axis('equal');

# Calculate covariance matrix
```

```
cov_mat(Y.T)
array([[ 4.94072998, -4.93536067],
       [-4.93536067,  5.99552455]])
```

Copy



This leads to the question of how to decompose the covariance matrix  $C$  into a rotation matrix  $R$  and a scaling matrix  $S$ .

## Eigen Decomposition of the Covariance Matrix

Eigen Decomposition is one connection between a linear transformation and the covariance matrix. An eigenvector is a vector whose direction remains unchanged when a linear transformation is applied to it. It can be expressed as

$$Av = \lambda v$$

where  $v$  is an eigenvector of  $A$  and  $\lambda$  is the corresponding eigenvalue. If we put all eigenvectors into the columns of a Matrix  $V$  and all eigenvalues as the entries of a diagonal matrix  $L$  we can write for our covariance matrix  $C$  the following equation

$$CV = VL$$

where the covariance matrix can be represented as

$$C = VLV^{-1}$$

which can be also obtained by [Singular Value Decomposition](#). The eigenvectors are unit vectors representing the direction of the largest variance of the data, while the eigenvalues represent the magnitude of this variance in the corresponding directions. This means  $V$  represents a rotation matrix and  $\sqrt{L}$  represents a scaling matrix. From this equation, we can represent the covariance matrix  $C$  as

$$C = RSSR^{-1}$$

where the rotation matrix  $R = V$  and the scaling matrix  $S = \sqrt{L}$ . From the previous linear transformation  $T = RS$  we can derive

$$C = RSSR^{-1} = TT^T$$

because  $T^T = (RS)^T = S^T R^T = SR^{-1}$  due to the properties  $R^{-1} = R^T$  since  $R$  is orthogonal and  $S = S^T$  since  $S$  is a diagonal matrix. This enables us to calculate the covariance matrix from a linear transformation. In order to calculate the linear transformation of the covariance matrix, one must calculate the eigenvectors and eigenvalues from the covariance matrix  $C$ . This can be done by calculating

$$T = V\sqrt{L}$$

where  $V$  is the previous matrix where the columns are the eigenvectors of  $C$  and  $L$  is the previous diagonal matrix consisting of the corresponding eigenvalues. The transformation matrix can be also computed by the [Cholesky decomposition](#) with  $Z = L^{-1}(X - \bar{X})$  where  $L$  is the Cholesky factor of  $C = LL^T$ .

We can see the basis vectors of the transformation matrix by showing each eigenvector  $v$  multiplied by  $\sigma = \sqrt{\lambda}$ . By multiplying  $\sigma$  with 3 we cover approximately 99.7% of the points according to the [three sigma rule](#) if we would draw an ellipse with the two basis vectors and count the points inside the ellipse.

```
C = cov_mat(Y.T)
eVe, eVa = np.linalg.eig(C)

plt.scatter(Y[:, 0], Y[:, 1])
for e, v in zip(eVe, eVa.T):
    plt.plot([0, 3*np.sqrt(e)*v[0]], [0, 3*np.sqrt(e)*v[1]], 'k-', lw=2)
plt.title('Transformed Data')
plt.axis('equal');
```



We can now get from the covariance the transformation matrix  $T$  and we can use the inverse of  $T$  to remove correlation (whiten) the data.

```
C = cov_mat(Y.T)

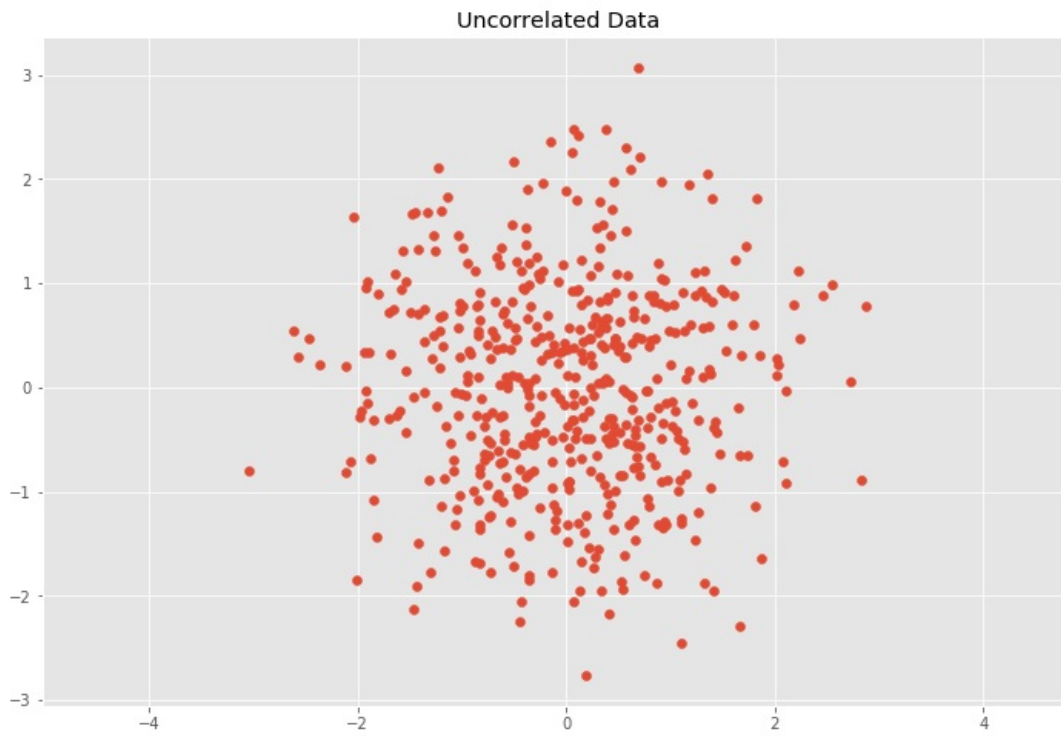
# Calculate eigenvalues
eVa, eVe = np.linalg.eig(C)

# Calculate transformation matrix from eigen decomposition
R, S = eVe, np.diag(np.sqrt(eVa))
T = R.dot(S).T

# Transform data with inverse transformation matrix T^-1
Z = Y.dot(np.linalg.inv(T))

plt.scatter(Z[:, 0], Z[:, 1])
plt.title('Uncorrelated Data')
plt.axis('equal');
```

# Covariance matrix of the uncorrelated data  
`cov_mat(Z.T)`  
`array([[ 1.00000000e+00, -1.24594167e-16],`  
 `[-1.24594167e-16, 1.00000000e+00]])`



An interesting use of the covariance matrix is in the [Mahalanobis distance](#), which is used when measuring multivariate distances with covariance. It does that by calculating the uncorrelated distance between a point  $x$  to a multivariate normal distribution with the following formula

$$D_M(x) = \sqrt{(x-\mu)^T C^{-1} (x-\mu)}$$

where  $\mu$  is the mean and  $C$  is the covariance of the multivariate normal distribution (the set of points assumed to be normal distributed). A derivation of the Mahalanobis distance with the use of the Cholesky decomposition can be found in this [article](#).